

Experiment 5: K-Means Clustering after Autoencoder Dimensionality Reduction

Imports and Configuration

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%cd ..
from src.models.autoencoder import Autoencoder
from src.models.kmeans import KMeans, KMeansInit
from src.utils.metrics import *
from src.utils.utils import *

RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)

c:\Users\ziade\OneDrive\Desktop\Term 7\ML\Lab 4\Assignment4-
Unsupervised-Clustering-Analysis
```

Data Loading and Preprocessing

```
X_scaled, y_true =
load_scale_data(data_path='./data/breast_cancer.csv')

print(f"Dataset shape: {X_scaled.shape}")
print(f"Labels shape: {y_true.shape}")

Dataset shape: (569, 30)
Labels shape: (569,)
```

Training and Evaluation Loop

```
bottleneck_dims = [2, 5, 10, 15, 20] #
ae_kms_results = []
all_loss_histories = {}

for b_dim in bottleneck_dims:
    print(f"Training Autoencoder with Bottleneck: {b_dim}...")

    # Initialize Autoencoder [cite: 18, 20, 24]
    ae = Autoencoder(
        input_dim=X_scaled.shape[1],
        encoding_dims=[64, 32], # Two hidden + bottleneck = 3 layers
        bottleneck_dim=b_dim,
        activation='relu',
```

```

        learning_rate=0.001,
        l2_lambda=0.01, # Regularization [cite: 24]
        epochs=150,
        batch_size=32 # Mini-batch GD [cite: 22]
    )

# Train and record loss [cite: 92]
ae.fit(X_scaled, verbose=False)
all_loss_histories[b_dim] = ae.train_loss_history

# 1. Dimensionality Reduction Quality [cite: 89, 90]
recon_error = ae.reconstruction_error(X_scaled)
X_latent = ae.encode(X_scaled)

# 2. Clustering in Latent Space [cite: 59]
km = KMeans(K=2, init=KMeansInit.KMEANS_PLUS_PLUS,
random_state=42)
km.fit(X_latent)
y_pred = km.clusters

# 3. Validation Metrics [cite: 75, 76, 83]
ae_kms_results.append({
    'Bottleneck': b_dim,
    'Reconstruction MSE': recon_error,
    'ARI': compute_ari(y_true, y_pred),
    'NMI': compute_nmi(y_true, y_pred),
    'Silhouette': compute_silhouette_score(X_latent, y_pred),
    'Purity': purity_score(y_true, y_pred),
    'DB_Index': compute_davies_bouldin(X_latent, y_pred)
})

```

results_df = pd.DataFrame(ae_kms_results)
results_df


Training Autoencoder with Bottleneck: 2...
Training Autoencoder with Bottleneck: 5...
Training Autoencoder with Bottleneck: 10...
Training Autoencoder with Bottleneck: 15...
Training Autoencoder with Bottleneck: 20...

	Bottleneck	Reconstruction MSE	ARI	NMI	Silhouette
Purity \					
0	2	0.999891	0.002817	0.000996	0.991565
0.627417					
1	5	0.828903	0.104615	0.080776	0.824460
0.687170					
2	10	0.796806	0.208906	0.309246	0.590737
0.731107					
3	15	0.472280	0.331896	0.330539	0.625165
0.796134					

4	20	0.376416	0.200915	0.154022	0.533427
0.736380					
DB_Index					
0	0.097722				
1	0.444631				
2	0.617133				
3	0.702903				
4	0.941282				

Training Loss Curves

```
plt.figure(figsize=(10, 6))
for b_dim, history in all_loss_histories.items():
    plt.plot(history, label=f'Bottleneck {b_dim}')
plt.title("Autoencoder Training Loss Curves (MSE + L2)")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```




Reconstruction Loss vs. Clustering Performance Analysis

```
fig, ax1 = plt.subplots(figsize=(10, 5))

ax1.set_xlabel('Bottleneck Dimension')
ax1.set_ylabel('Reconstruction MSE', color='red')
ax1.plot(results_df['Bottleneck'], results_df['Reconstruction MSE'],
          color='red', marker='o', label='MSE')
ax1.tick_params(axis='y', labelcolor='red')

ax2 = ax1.twinx()
ax2.set_ylabel('Adjusted Rand Index (ARI)', color='blue')
ax2.plot(results_df['Bottleneck'], results_df['ARI'], color='blue',
          marker='s', label='ARI')
ax2.tick_params(axis='y', labelcolor='blue')

plt.title("Reconstruction Loss vs. Clustering Quality")
plt.show()
```




2D Projection and Clusters

```
# Overlay clusters and true labels
plt.figure(figsize=(8, 6))
sns.scatterplot(
    x=X_2d[:, 0],
    y=X_2d[:, 1],
    hue=km_2d.clusters,
    style=y_true, # Shape encodes actual labels
    palette='viridis'
)
```

```

plt.title("2D Latent Space Projection (Bottleneck=2) with True Labels Overlay")
plt.xlabel("Latent Dim 1")
plt.ylabel("Latent Dim 2")
plt.legend(title="Cluster / True Label")
plt.show()

```



```

# Find best bottleneck by ARI
best_bottleneck = results_df.loc[results_df['ARI'].idxmax()]
['Bottleneck']
ae_best = Autoencoder(X_scaled.shape[1], [64, 32],
int(best_bottleneck))
ae_best.fit(X_scaled, verbose=False)
X_best = ae_best.encode(X_scaled)
km_best = KMeans(K=2, init=KMeansInit.KMEANS_PLUS_PLUS,
random_state=42)
km_best.fit(X_best)
y_pred_best = km_best.clusters

# Align clusters to labels

```

```

y_aligned = align_clusters_with_labels(y_true, y_pred_best)

cm_ae = compute_confusion_matrix(y_true, y_aligned)
plt.figure(figsize=(5, 4))
sns.heatmap(cm_ae, annot=True, fmt='d', cmap='Oranges',
            xticklabels=['Benign (0)', 'Malignant (1)'],
            yticklabels=['Benign (0)', 'Malignant (1)'])
plt.title(f"Confusion Matrix: Autoencoder + K-Means  
(Bottleneck={best_bottleneck})")
plt.ylabel("True Label")
plt.xlabel("Predicted Label")
plt.show()

```

Confusion Matrix: Autoencoder + K-Means (Bottleneck=15.0)

